

Powering the Future: Energy Harvesting and Intermittent Computing in IoT

Anton Odén

Dept. of Maths and Computer Science

Karlstad University

651 88 KARLSTAD, Sweden

anton.oden@outlook.com

Abstract—

CONTENTS

A. Introduction

B. Energy harvesting: A new power paradigm

Energy harvesting refers to the process of capturing and converting ambient energy from natural or artificial sources into usable electrical power. Common sources of harvested energy include solar radiation, thermal gradient, mechanical vibration, biochemical reactions and radio frequencies [Energyharvest1]. Unlike traditional power supplies, energy harvesting systems enable autonomous operation in environments where wired connection or battery replacements are impractical. This technique is increasingly vital for IoT devices, embedded systems with batter-free computing where maintaining continuous power supply is challenging or that cutting the cost of battery is the next step in making IoT devices even more cheap.

One promising and intriguing technology in energy harvesting is RF (radio frequency) energy harvesting, which utilizes electromagnetic waves from sources such as cell towers, WiFi signals and RFID readers to generate power. The WISP (Wireless identification and sensing platform) project exemplifies this approach, demonstrating how RF energy can support battery-free computation and sensing. WISP operates using passive RFID technology, converting radio waves into energy enough for tasks like data processing and sensor driven applications. This enables devices to function in environments with consistent RF energy exposure, paving the way for battery-free IoT solutions.

Despite advantages of IoT devices able to compute and communicate without traditional battery. Energy harvesting does have a challenge with ambient energy fluctuation leading to unpredictable power levels. This leads to systems required to adopt intermittent computing strategies such as e.g. Alpaca, Ratchet and Chinchilla further discussed in this article.

C. Understanding intermittent computing

Intermittent computing is designed for systems that experience frequent power interruptions, often seen in energy-




Fig. 1. Different energy harvesting areas showing given energy from [SurveySecurity]

harvesting devices. Unlike traditional computing, which assumes continuous power availability, intermittent systems must execute tasks in small steps, ensuring progress even during power failures. Counter to traditional computing that could be considered to be continuous. Intermittent computing is able to handle planned or unplanned breaks in computation. The processor in itself does not have a problem with powerbreaks but the memory used in computation to write and read data is more or less stored in volatile memory, for example RAM, Cache, CPU registers. Meaning that it doesn't uphold data during power loss. Memory that does keep data during power loss is non-volatile and could be for example FRAM, Flash, EEPROM. Traditional computing then fails under unstable power as memory used by computation is lost and there is no checkpoints to know what state the computation and memory was in close to power loss so the computation results can't be trusted to be correct if it wouldn't restart from beginning. Energy harvesting dependent systems have to operate intermittently as the energy harvested can't be assumed to be constant, but instead is losing power frequently and unexpectedly. Since power availability differs in energy-harvesting systems, computing strategies must adapt. Some key approaches include, **checkpointing** that regularly saves system state so computation can resume at last checkpoint (e.g. Chinchilla

and Ratchet). **Task-based execution** where program is broken into small self-contained tasks that restarts after power loss (Alpaca). **Non-volatile memory** usage, storing of execution data in FRAM or Flash memory to preserve progress. **Energy aware scheduling** that could dynamically adjust workload based on realtime power levels, ensuring tasks execute only when sufficient energy is available.

D. Frameworks for intermittent computing

Intermittent computing requires specialized compilation techniques to ensure programs can function despite unpredictable power failures. Unlike traditional compilation, which assumes continuous execution, intermittent computing frameworks modify code execution behavior so systems can resume operations seamlessly after a power outage. These frameworks aim to prevent data loss when power fails, enable efficient recovery without manual intervention and optimize execution for low-power environments.

1) *Why compilation matters for intermittent systems:* Since energy-harvesting devices operate with instable power, normal software execution fails because variables, loop states and memory data disappear when power is lost. Specialized compilation frameworks enhance program by automatically adding checkpoints to preserve state (Ratchet) or adjusts checkpoint frequency dynamically based on power availability (Chinchilla). Also iterated computation statements could in compilation be flagged as non-idempotent and temporary variables are added by compilers (Alpaca).

2) *Ratchet: Automatic checkpointing without hardware support:* Ratchet provides compiler-based automatic checkpointing which makes it possible to use without hardware support, meaning standard microcontrollers could be used. Ratchet inserts checkpoints in compilation without requiring programmer to modify code manually.

Ratchet transform execution state into safe recoverable snapshots, ensuring program resume smoothly after power loss. It ensures redundancy is minimized to prevent unnecessary recomputation.

3) *Chinchilla: adaptive checkpoints for efficient computing:* Chinchilla introduces dynamic checkpointing that adjust checkpoint frequency based on available energy to maximize efficiency. This ensures minimal overhead by choosing only necessary checkpoints instead of static as in Ratchet.

4) *Alpaca: Task-based programs for software resilience:* In Alpaca more work is on the programmer to divide programs into chunks of tasks. The programmer needs to be aware of the amount of energy that is possible at a minimum to compute with so a task is not too large and thereby not able to finish. As the program overhead is handed to the programmer there is more room for designing efficient code adapted to the device environment which could differ largely between devices. Overhead is limited to every individual task write to non-volatile memory at the end of a task. Second feature-leg for Alpaca is privatization that guarantees that volatile or non-volatile memory access by a task remains consistent, regardless of power failure. Alpaca divides data

into task-shared and task-local data. Task-shared variables are in the global scope and are allocated in non-volatile memory and once a task writes a value to a task-shared variable, that same task or another task may later read the value. In Alpaca the compiler backtracks write and reads to all task-shared variables to decide if tasks need to privatize usage of variables, meaning that a local variable is created in tasks and in the end of task a commit to task-shared variable is done before transitioning to next task. Task-local variables are scoped only to a single task and must be initialized by that task and are allocated in volatile memory. Committing of data from volatile to non-volatile memory is done in two phases and different index set to true or false is set to signal what phase the task was in if it's been power interrupted. First pre-commit creates a table in non-volatile memory with all values to change in the global scope. If interrupted during this phase the whole task needs to be computed from beginning. After pre-commits done index bit commit, ready is set to true and Alpaca runs until during next boot known to job. Shared variables are changed in volatile memory, same as with pre-commit but with a commit, ready flag is changed after commit phase is done. Shared variables for Alpaca to guarantee consistency of data.

E. Future directions and applications

Hello Hello Hello

F. Conclusion

In this article we have discussed intermittent computing, energy harvesting of radio waves and how IoT could potentially power much itself by sheer device volume.

Intermittent computing adds a layer of consideration for the programmer and designer of device. We have highlighted Ratchet, Chinchilla and Alpaca that are three approaches to tackle the problem of keeping track of program state with power interruptions that devices being reliant on energy harvesting only has. But of course power interruptions could happen otherwise also so the concepts of adding checkpoints is applicable in a broader sense although this article focuses on smaller IoT devices.